

Who told the agent it **could do that?**

Accountability & auditability — making every agent action attributable, replayable, and something you can actually defend.

[Swipe →](#)

SIX WEEKS LATER

“Why did it remove those licenses on the 14th — and who signed off?”

A fair question from the customer’s security team. We went to the logs.

```
> Task completed successfully. 42 actions taken.
```

We knew **that** it acted. We couldn’t reconstruct **why**, what it saw, or who was accountable.

The question that ends pilots isn't "does it work?"

It's "can you prove what it did?"

Autonomy erases the human paper trail. A person clicked the button, their name was on it. When the **agent** clicks it, that trail evaporates — unless you rebuild it on purpose.

Four properties.

- **Attributable** — agent + version + plan + the identity that authorized it.
- **Replayable** — the exact context it had at decision time. Not a summary.
- **Explainable** — the reasoning, in a form a human can review.
- **Owned** — a named team is accountable. The model is never the owner.

Log **decisions**, not just outcomes.

Most systems log the effect: “*license removed, user 4412.*” That’s necessary and completely insufficient.

The old, proven instinct: **event sourcing**.
Don’t store only the final state — store the sequence of decisions that produced it.
Agents need the same discipline, one layer up.

Capture the whole decision.

Trigger

what initiated this step

Context

the inputs & tool results it actually had

Intent

what it was trying to accomplish

Action + Auth

the operation + who/what permitted it

Result

outcome, and any compensation

Now “why did it do X on the 14th?” is a **query**, not a forensics project.

The audit trail is a **security control**.

Indirect prompt injection can turn a helpful agent into a confused deputy — content it reads carries hidden instructions, and it acts with its own privileges.

When that happens at scale, a decision-level trace is the only thing between you and **“we have no idea what it touched.”**

An accountability checklist — and a trace schema.

Is every action **attributable**? Can you **replay**
the context? Is the reasoning **explainable**?
Does every agent have a named **owner**?

Full post + checklist + copy-paste trace schema —
link in the comments.

Part 2 = reliability. Part 3 = evaluation. Next: Part 5
— Context & Memory.